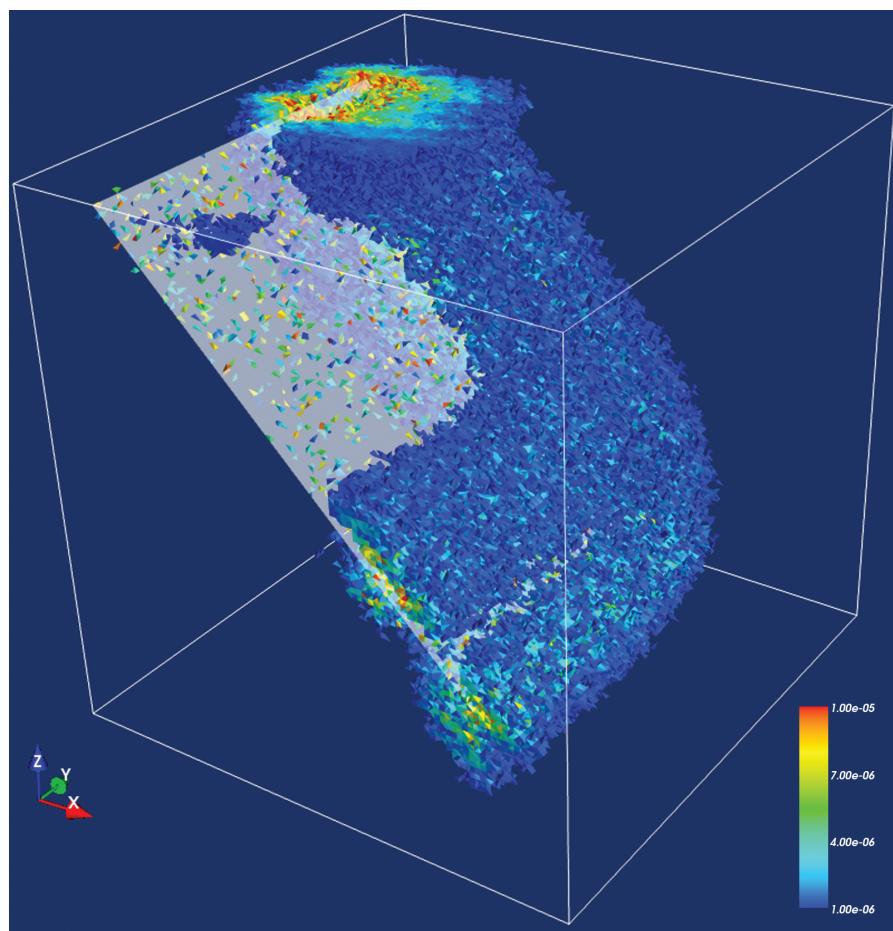


Cigma

User Manual
Version 0.9.1



www.geodynamics.org

Luis Armendariz
Susan Kientz

Cigma

© California Institute of Technology
Version 0.9.1

March 6, 2008

Contents

1	Introduction	7
1.1	About Cigma	7
1.2	Citation	7
1.3	Support	8
2	Installation and Getting Help	9
2.1	Getting Help	9
2.2	Installing from Source	9
2.2.1	HDF5	9
2.2.2	VTK	9
2.2.3	NumPy	10
2.2.4	PyTables	10
2.2.5	FIAT	10
2.2.6	HDFView (optional)	10
3	Error Analysis	13
3.1	Introduction	13
3.2	Distance Measures	13
3.3	Global Error	14
3.4	Comparing Residuals	14
3.5	Input and Output	14
3.6	Mesh	15
3.7	Fields	15
3.8	Elements	15
4	Running Cigma	17
4.1	Listing Data	17
4.2	Comparing Two Fields	18
4.2.1	Specifying a Mesh	18
4.2.2	Specifying a Quadrature Rule	18
4.2.3	Comparing against Known Values	18
4.2.4	Comparing against a Known Function	19
5	Examples	21
5.1	Convergence	21
5.2	Different Element Types	21
5.3	Comparing Two Codes	22
5.4	Analytic Comparison	22

A File Formats	23
A.1 Input File Format	23
A.1.1 HDF5	23
A.1.2 VTK	24
A.1.3 Text	24
A.2 Output File Format	25
B License	27

List of Figures

- | | |
|---|----|
| 2.1 With HDFView, you can view the internal file hierarchy in a tree structure, add new datasets,
and modify or delete existing datasets | 11 |
|---|----|

Chapter 1

Introduction

1.1 About Cigma

The CIG Model Analyzer (Cigma) consists of a suite of tools for comparing numerical models. CIG has developed Cigma in response to demand from the short-term tectonics community for a simple tool that can perform rigorous error analysis on their finite element codes. The long-term goal for Cigma, however, is for it to be used for nearly all geodynamics modeling codes.

The current version of Cigma is intended for the calculation of L_2 residuals for finite element models. This software was designed for performing the following three main tasks: (1) error analysis, (2) benchmarking, and (3) code verification.

In error analysis, Cigma can calculate a global measure of the error between two solutions by performing an integration over a discretized version of the common domain. This comparison can take place even when the underlying discretizations do not overlap.

In benchmarking, Cigma can help the geodynamics community agree on a standard solution to specific problems by facilitating the process of comparing different numerical codes against each other.

Lastly, as an automated tool, Cigma can help application developers create regression tests to ensure that software changes do not affect the consistency of the results.

At its core, Cigma draws from a variety of libraries. The FInite element Automatic Tabulator (FIAT) (www.fenics.org/wiki/FIAT) Python Library supports generation of arbitrary order instances of the Lagrange elements on lines, triangles, and tetrahedra. It can also generate order instances of Jacobi-type quadrature rules on the same element shapes. The Approximate Nearest Neighbor (ANN) (www.cs.umd.edu/~mount/ANN/) Library, written in C++, supports data structures and algorithms for both exact and nearest-neighbor searching in arbitrarily high dimensions.

Both of these libraries particularly extend and generalize Cigma's functionality so it can handle other types of elements, and provide the ability to compare vector fields.

1.2 Citation

Computational Infrastructure for Geodynamics (CIG) is making this source code available to you in the hope that the software will enhance your research in geophysics. This is a brand-new code and at present no papers are published or at press for use as citations other than this manual, which is cited as follows:

- Armendariz, L., and S. Kientz. *Cigma User Manual*. Pasadena, CA: Computational Infrastructure of Geodynamics, 2007. URL: geodynamics.org/cig/software/cs/cigma.pdf

CIG requests that in your oral presentations and in your papers that you indicate your use of this code and acknowledge the author of the code and CIG (geodynamics.org).

1.3 Support

Cigma development is supported by a grant from the National Science Foundation to CIG, managed by the California Institute of Technology. The code is being released under the GNU General Public License.

Chapter 2

Installation and Getting Help

2.1 Getting Help

For help, send an e-mail to the CIG Computational Science Mailing List (cig-cs@geodynamics.org). You can subscribe to the `cig-cs` mailing list and view archived discussions at the CIG Mail Lists web page (geodynamics.org/cig/lists). If you encounter any bugs or have problems installing Cigma, please submit a report to the CIG Bug Tracker (geodynamics.org/bugs).

2.2 Installating from Source

To install Cigma, download the source package from the CIG Cigma web page (geodynamics.org/cig/software/packages/cs/cigma). You will need the GNU C++ compiler for this step. The HDF5 and VTK libraries, described later in this section, are also required.

After installing the necessary dependencies, simply unpack the source tar file and issue the following commands

```
$ tar xvfz cigma-0.9.1.tar.gz
$ cd cigma-0.9.1
$ ./configure --with-hdf5=$HDF5_PREFIX --with-vtk=$VTK_PREFIX
$ make
$ sudo make install
```

2.2.1 HDF5

The Hierarchical Data Format (HDF5) library is available for download from The HDF Group (hdfgroup.org/HDF5). Binaries can be obtained at hdfgroup.org/HDF5/release/obtain5.html (hdfgroup.org/HDF5/release/obtain5.html). To install from source, download the latest stable 1.6 version of this library (currently 1.6.5) and issue the following commands

```
$ tar xvfz hdf5-1.6.5
$ cd hdf5-1.6.5
$ ./configure
$ make
$ sudo make install
```

2.2.2 VTK

The Visualization Toolkit (VTK) library is a popular open source graphics library for scientific visualizations. We recommend that you obtain the binaries from a package manager, making sure to install the associated development headers along with the library. The source for the VTK library is available from Kitware,

Inc. (www.vtk.org/get-software.php). If you decide to compile VTK from source, you will also need the CMake build environment, available from CMake (cmake.org).

After downloading the source package, you can issue the following commands

```
$ tar xvfz vtk-5.0.4.tar.gz
$ cd VTK
$ mkdir build
$ cd build
$ ccmake ..
$ make
$ sudo make install
```

2.2.3 NumPy

NumPy is a Python module which adds support for large multi-dimensional arrays and matrices, and is available for download from the NumPy Home Page (numpy.scipy.org). To install this extension module from source, download it and issue the following command in the NumPy source directory

```
$ sudo python setup.py install
```

To verify the installation, the command ‘`import numpy`’ should run successfully from within your standard Python interpreter shell.

2.2.4 PyTables

PyTables is a Python extension module that builds on top of the HDF5 library. It provides a convenient scripting interface to manipulate HDF5 files, which can be used to manipulate the input/output files created by Cigma. PyTables is available for download from PyTables (www.pytables.org).

To install this extension from source, download the latest stable version (currently 2.0.2) and issue the following commands

```
$ tar xvfz pytables-2.0.2
$ cd pytables-2.0.2
$ sudo python setup.py install
```

To verify the installation, the command ‘`import tables`’ should run successfully from within your standard Python interpreter shell.

2.2.5 FIAT

The Finite Element Automatic Tabulator (FIAT) Python module supports generation of arbitrary order instances of the Lagrange elements, as well as arbitrary order instances of the Jacobi-type quadrature rules on the same element shapes. It is available from The FEniCS Project (www.fenics.org/wiki/FIAT).

To install this module, download it and issue the following command inside the FIAT source directory:

```
$ sudo python setup.py install
```

You should now be able to run ‘`import FIAT`’ from within your standard Python interpreter shell.

2.2.6 HDFView (optional)

NCSA HDFView is a graphical user interface tool for accessing data in your HDF5 files. You can use it for viewing the internal file hierarchy in a tree structure, adding new datasets, and modifying or deleting existing datasets, or altering metadata on groups and datasets. You can download it from the HDFView home page (hdf.ncsa.uiuc.edu/hdf-java-html/hdfview).

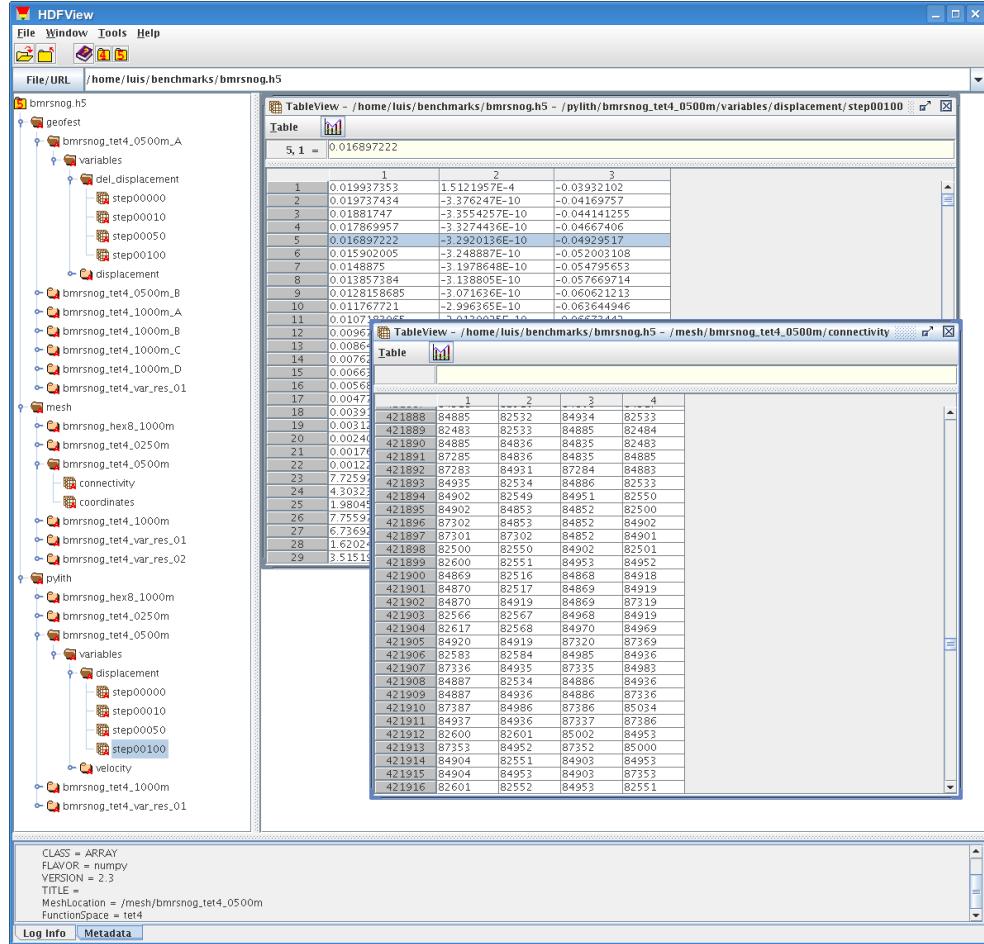


Figure 2.1: With HDFView, you can view the internal file hierarchy in a tree structure, add new datasets, and modify or delete existing datasets

Chapter 3

Error Analysis

3.1 Introduction

When studying differential equations that represent physical systems we often obtain solutions by using a variety of techniques. Solving these equations using classical analytical methods is almost impossible, so we often resort to numerical algorithms such as the finite element method.

For assessing the quality of the solution to our equations, it is important to quantify the error in our numerical solutions. To calculate this error, we will have to compare tentative solutions against each other through the use of a distance measure between two functions.

3.2 Distance Measures

The simplest possible quantitative measure of the difference between two distinct fields you can make consists of taking the pointwise difference of both fields at a common set of points. While no finite sample of points can perfectly represent a continuum of values, valuable information can be inferred from a statistics analysis of the resulting set of differences.

Perhaps a more useful distance measure can be obtained by using the L_2 norm, defined by the following integral

$$\varepsilon = \|u - v\|_{L_2} = \sqrt{\int_{\Omega} \|u(\vec{x}) - v(\vec{x})\|^2 d\vec{x}}$$

This gives us a single global estimate ε representing the distance between the two fields $u(\vec{x})$ and $v(\vec{x})$. Alternatively, you may think of this as the size, or norm, of the residual field $\rho(\vec{x}) = u(\vec{x}) - v(\vec{x})$.

If we discretize the domain Ω into finite elements Ω_e , the above integral can be broken up into a sum over local contributions on each element. For efficiency, each contribution can be integrated over a reference element $\hat{\Omega}_e$ defined on a standard coordinate system.

$$\begin{aligned}\varepsilon^2 &= \sum_{e=1}^{nel} \varepsilon_e^2 \\ &= \sum_{e=1}^{nel} \int_{\Omega_e} \|u(\vec{x}) - v(\vec{x})\|^2 d\vec{x} \\ &= \sum_{e=1}^{nel} \int_{\hat{\Omega}_e} \|u(\vec{\xi}) - v(\vec{\xi})\|^2 J(\vec{\xi}) d\vec{\xi}\end{aligned}$$

In general, we won't be able to integrate each local contribution exactly since the two fields u and v may have a representation that is incompatible with the local domain Ω_e . However, we can approximate each ε_e^2 by applying an appropriate quadrature rule with a tolerable truncation error [2].

3.3 Global Error

Assuming we apply the same quadrature rule, with weights w_q and integration points $\vec{\xi}_q$, on every element,

$$\begin{aligned}\varepsilon_e^2 &= \sum_{q=1}^{nq} w_q \|\hat{\rho}(\vec{x}_q)\|^2 \\ &= \sum_{q=1}^{nq} w_q \|u(\vec{\xi}_q) - v(\vec{\xi}_q)\|^2 J(\vec{\xi}_q)\end{aligned}$$

thus we arrive at the final form

$$\varepsilon = \sqrt{\sum_{e=1}^{nel} \sum_{q=1}^{nq} w_q \|u(\vec{\xi}_q) - v(\vec{\xi}_q)\|^2 J(\vec{\xi}_q)}$$

In calculating the norm of the residual field ρ , Cigma will output each of the local contributions ε_e^2 which by definition are scalar-valued cell quantities over each of their corresponding elements Ω_e .

3.4 Comparing Residuals

For comparing errors of different solutions, you can normalize the global error by the norm of the exact solution. This normalized error is given by

$$\begin{aligned}\varepsilon_{rel} &= \frac{\|u - u^h\|_{L_2}}{\|u\|_{L_2}} \\ &= \frac{\int_{\Omega} \|u(\vec{x}) - u^h(\vec{x})\|^2 d\vec{x}}{\int_{\Omega} \|u(\vec{x})\|^2 d\vec{x}}\end{aligned}$$

This normalized error can be interpreted as the average error in the physical quantity being evaluated, so that a value of 0.01 corresponds to a 1% averaged error. The norm in the denominator can be computed in Cigma by comparing against the built-in zero function (see Section 4.2.4).

Even if the exact solution is not currently known, this normalized error may be used to test the accuracy between two or more numerical solutions, defined on successively refined meshes.

3.5 Input and Output

The underlying data storage format for Cigma is the HDF5 format, due to its flexibility for storing and organizing large amounts of data.

The Hierarchical Data Format (HDF) is a portable file format developed at the National Center for Supercomputing Applications (NCSA) (hdf.ncsa.uiuc.edu/HDF5). It is designed for storing, retrieving, analyzing, visualizing, and converting scientific data. The current and most popular version is HDF5, which stores multi-dimensional arrays together with ancillary data in a portable self-describing format. It uses a hierarchical structure that provides application programmers with a host of options for organizing how data is stored in HDF5 files.

Using HDF5 datasets in Cigma allows us to avoid having to convert between too many distinct formats. Moreover, due to the amount of disk I/O, large finite element meshes can be handled more efficiently in binary format.

Cigma also supports the popular VTK format for certain kinds of input. A simple text format is also supported to aid in debugging. As described in Chapter 4, you can easily examine the structure of an input file by using the `cigma list` command, which will simply reveal the names and dimensions of all datasets inside the specified file.

3.6 Mesh

In Cigma, we define a finite element mesh simply by coordinates, (x_n, y_n, z_n) of its degrees of freedom, and the connectivity relations $\Omega_e = \{n_1, n_2, \dots\}$ among them which define each individual element in the corresponding discretization.

Since the evaluation points on two distinct meshes will (TODO – will do what?), it is very important to index the location of each element into spatial data structure.

3.7 Fields

A field is a function which assigns a physical quantity to every point in space. This quantity may correspond to a scalar, a vector, or even a tensor. A finite element approximation to an unknown field $\phi(\vec{x})$ is simply the weighed sum over a fixed set of localized shape functions $\phi_n(\vec{x})$.

$$\phi(\vec{x}) = \sum_{n=1}^N d_n \phi_n(\vec{x})$$

Naturally, once the weights d_n , or degrees of freedom as they are also called, are known to us, we can evaluate $\phi(\vec{x})$ at any point \vec{x} we desire. Thus, in Cigma a field is represented simply by a list of degrees of freedom d_n , which may be a scalar, vector or tensor quantity, depending on the nature of ϕ .

3.8 Elements

The most common element types used in the finite element method are given by the following shape functions, defined on their respective domains,

Function Space tet4

$$\begin{aligned} TN_a &= \frac{1}{2}(-1 - x - y - z) \\ TN_b &= \frac{1}{2}(1 + x) \\ TN_c &= \frac{1}{2}(1 + y) \\ TN_d &= \frac{1}{2}(1 + z) \end{aligned}$$

Function Space hex8

$$\begin{aligned} HN_a &= \frac{1}{8}(1 - x)(1 - y)(1 - z) \\ HN_b &= \frac{1}{8}(1 + x)(1 - y)(1 - z) \\ HN_c &= \frac{1}{8}(1 - x)(1 + y)(1 - z) \\ HN_d &= \frac{1}{8}(1 + x)(1 + y)(1 - z) \\ HN_e &= \frac{1}{8}(1 - x)(1 - y)(1 + z) \\ HN_f &= \frac{1}{8}(1 + x)(1 - y)(1 + z) \\ HN_g &= \frac{1}{8}(1 - x)(1 + y)(1 + z) \\ HN_h &= \frac{1}{8}(1 + x)(1 + y)(1 + z) \end{aligned}$$

These are built-in elements in Cigma, and you may refer to them by name in the FunctionSpace metadata attribute on your datasets.

Higher order elements can also be used by specifying an explicit tabulation matrix representing the values of your shape functions on the corresponding set of quadrature points.

Chapter 4

Running Cigma

Cigma is designed to be scriptable. Thus, all operations can be specified as command line arguments given to a single executable called `cigma`. A list of available commands can be obtained by typing

```
cigma --help
```

Further usage information can also be obtained with

```
cigma help <command>
```

Cigma is used for obtaining error estimates between arbitrary fields, so naturally its primary operation involves the `compare` command. You will need to provide two datasets describing each of the two fields, along with an integration rule and a mesh over which to integrate, although these last two will have reasonable defaults if they are not specified.

Specifying the complete path to a dataset consists of the special form `filepath:dataset`, a colon-delimited pair of file path and dataset path.

4.1 Listing Data

Because Cigma relies on the ability of the user to specify dataset paths, we have provided a command called `list` for viewing the structure of an input file. Its usage is very simple.

To view the structure of an HDF5 file:

```
$ cigma list file.h5
/mesh/coordinates          Dataset {119827, 3}
/mesh/connectivity          Dataset {661929, 4}
/vars/displacement/step0    Dataset {119827, 3}
/residuals/comparison0     Dataset {661929, 1}
```

If VTK support is enabled, you can view the structure of a VTK file with:

```
$ cigma list file.vtk
Reading file.vtk
Points = 119827
Cells = 661929
PointDataArray[0] = displacements_t0 (119827 x 3)
```

4.2 Comparing Two Fields

Comparing two arbitrary finite element fields can be accomplished with the `cigma compare` command-line utility. By default, the comparison will take place over the elements in the mesh of the first field. Finally, the square of each of the local residual values are written to the specified VTK output file as cell-based scalars.

A basic comparison can be as simple as specifying the following arguments:

```
cigma compare --output=squared-residuals.vtk
  --first=field1.h5:/field1/stepN
  --second=field2.h5:/field2/stepN
```

4.2.1 Specifying a Mesh

To override the mesh used in the integration, you can specify an extra argument providing the location of the mesh:

```
cigma compare [...]
  --mesh=mesh.h5:/model/mesh/
```

You can also specify the coordinates and connectivity arrays separately, in case they reside in separate files.

```
cigma compare [...]
  --mesh-coordinates=file1.h5:/model/mesh/coordinates
  --mesh-connectivity=file2.h5:/model/mesh/connectivity
```

4.2.2 Specifying a Quadrature Rule

To specify a quadrature rule, you will have to provide the quadrature weights and points in the appropriate reference element. This can be done with the following additional argument:

```
cigma compare [...]
  --rule=quadrature-rules.h5:/path/to/rule
```

You may also specify the location of the points and weights separately:

```
cigma compare [...]
  --rule-points=file.h5:/path/to/rule/points
  --rule-weights=file.h5:/path/to/rule/weights
```

4.2.3 Comparing against Known Values

A finite element description might not always be available for one of the fields. However, you can break the comparison into several steps if you have a means to compute that field on any of the required points.

First, extract the global coordinates of the integration points. This will result in an explicit list of points over which to evaluate your field.

```
cigma extract --mesh=field1.h5:/model/mesh/
  --output=points.h5:/points
```

Now you can evaluate your function at the designated points. You can provide the path to the explicit set of values with

```
cigma compare --output=squared-residuals.vtk
  --first=field1.h5:/field1/stepN
  --second=values.h5:/stepN_values
```

4.2.4 Comparing against a Known Function

If one of your fields is easily described by an analytic expression, then you also have the option to compile your analytic function into Cigma, which will enable the `compare` command to reference your function by name:

```
cigma compare --output=squared-residuals.vtk  
    --first=field1.h5:/vars/displacement/step0  
    --second=disloc3d
```

You may also interact with your analytic function by using the `cigma eval` command, and obtain a set of values which may then be passed back to the `compare` command.

```
cigma eval --function=disloc3d  
    --points=points.h5:/points  
    --values=values.h5:/disloc3d_values  
cigma compare --output=squared-residuals.vtk  
    --first=field1.h5:/vars/displacement/step0  
    --second=values.h5:/disloc3d_values
```


Chapter 5

Examples

When the exact solution of your equations is known, the norm of the error for a finite element solution is easily computed through the procedure given in Chapter 3.

Another way of checking the accuracy of a solution is by rerunning the same problem with a finer mesh. Recall that in the finite element method, the weight functions and trial solutions are chosen so that the finite element method is convergent. In other words, as the element size h decreases, the solution tends to the correct solution. Thus, if the difference between the coarse and fine mesh solutions is small, we can conclude that the coarse mesh solution is quite accurate. However, if a solution changes noticeably with refinement of the mesh, the coarse mesh solution is inaccurate, and even the finer mesh may be inadequate as well.

5.1 Convergence

To demonstrate the convergence in the FEM method, in this example we will compare the error incurred in representing a continuous function by finite elements.

You can change the function by modifying the `TestFunction.cpp` source file. Once you recompile `cigma`, you will be able to refer to it by the name `test` when using the `cigma eval` command.

In the test directory you will find a series of successively refined meshes that you can use for testing. The original mesh files were created with the Gmsh, a three-dimensional finite element mesh generator, but you will be accessing the HDF5.

In the figure (TODO: regenerate plot), we show the log of the L_2 error norm as a function of the log of the element size h .

As can be seen from these results, the log of the error varies linearly with element size and the slope depends on the order of the element. Denoting the slope by α , then the error in the function can be expressed as

$$\log \varepsilon = \log C + \alpha \log h$$

where $\log C$ is an arbitrary constant, the y-intercept of the curve. The slope α is the rate of convergence of the element. We can rewrite this as

$$\varepsilon = Ch^\alpha$$

5.2 Different Element Types

For this example we will use sample datasets from the strike-slip benchmark case defined by the CIG Short-Term Tectonics working group. In this benchmark problem, we solve for the viscoelastic relaxation of stresses from a single finite earthquake while ignoring gravity. The problem is defined on a cube domain with sides of 24 km consisting of two layers of different material types. The top layer of the cube is nearly elastic, while the bottom layer is viscoelastic.

5.3 Comparing Two Codes

In this example, we will compare various codes

5.4 Analytic Comparison

(TODO: ask brad for his okada soln HDF5 file and link to it here)

When the analytic solution is known, we may use it directly when specifying a field argument to the `cigma compare` command, although this will require you to register your function as a built-in field. To demonstrate this capability, the current version of Cigma includes `disloc3d` as a built-in field.

(TODO: rerun these commands)

```
$cigma eval --function=disloc3d --output=tet4_1000m_qpts_displacement
```

Alternatively, you can also extract the required list of points on which you need to evaluate the exact solution. We can do this for the Pylith mesh in the previous section with the command

```
$ cigma extract
--mesh=strikeSlip.h5:/tet4_1000m/mesh/
--output=tet4_1000m_qpts.h5:/tet4_1000m_qpts
```

You can now take the list of points in a global coordinate system, and apply your function. For demonstration purposes, we can use the `disloc3d` built-in function. The comparison can then take place

```
$ cigma compare
--first=strikeSlip.h5:/tet4_1000m/displacement/step0000
--second=okada_solution.h5:/tet4_1000m_qpts_displacement
```

Here, we have taken the

Appendix A

File Formats

To differentiate between these formats, you will need to provide an appropriate extension for your input and output files. The supported extensions are currently `.h5` for HDF5, `.vtk` for VTK files, and `.txt` for simple text files.

A.1 Input File Format

There are about six different objects you will want to use as input to Cigma, all of which can be represented as two-dimensional datasets. These are the

- mesh coordinates
- mesh connectivity
- field coefficients (degrees of freedom)
- quadrature rule points
- quadrature rule weights
- shape function tabulation matrix

A.1.1 HDF5

HDF5 files are organized in a hierarchical structure, similar to a UNIX file system. Two types of primary objects, groups and datasets, are stored in this structure. A group contains instances of zero or more groups or datasets, while a dataset stores a multi-dimensional array of data elements. Both kinds of objects are accompanied by supporting metadata known as attributes.

A dataset is physically stored in two parts: a header and a data array. The header contains miscellaneous metadata describing the dataset as well as information that is needed to interpret the array portion of the dataset. Essentially, it includes the name, datatype, dataspace, and storage layout of the dataset. The name is a text string identifying the dataset. The datatype describes the type of the data array elements. The dataspace defines the dimensionality of the dataset, i.e., the size and shape of the multi-dimensional array.

In Cigma you always provide an explicit path to every dataset, so you have a fair amount of flexibility for how you organize your datasets inside HDF5 files. For example, a typical Cigma HDF5 file could have the following structure.

```
model.h5
\__ model
    \__ mesh
        \__ coordinates      [nno x nsd]
        \__ connectivity     [nel x ndof]
```

```

\_\_ variables
| \_\_ temperature
| | \_\_ step00000 [nno x 1]
| | \_\_ step00010 [nno x 1]
| | \...
| \_\_ displacement
| | \_\_ step00000 [nno x 3]
| | \_\_ step00010 [nno x 3]
| | \...
| \_\_ velocity
| | \_\_ step00000 [nno x 3]
| | \_\_ step00010 [nno x 3]
| |
\...

```

Generally, Cigma will only require you to specify the path to a specific field dataset. If a small amount of metadata is present in your field dataset, the rest of the required information, such as which mesh and finite elements to use, will be deduced from that metadata.

MeshID an identifier assigned for use in linking child datasets to their parent mesh.

MeshLocation points to the HDF5 group which contains the appropriate coordinates and connectivity datasets.

FunctionSpace string identifier to determine which shape functions to use for interpolating values inside the element (e.g., tet4, hex8, quad4, tri3, ...).

DatasetType string identifier for classifying the type of data contained in the dataset (e.g., points, connectivity, degrees of freedom, quadrature rules, global quadrature points, global field values).

A.1.2 VTK

The current version of Cigma only supports VTK unstructured grid datasets of a single element type. You can still compare various unstructured grids with different element types to each other.

Note that while you typically provide a path (or name) for every dataset, this is not necessary when specifying a VTK mesh, since this data is taken from the special Points and Cells arrays, which you cannot rename. However, you will still need to provide a name when referring to the shape function coefficients, which are assumed to be stored as Point Data in the input VTK file.

For more information, you may want to refer to Visualization ToolKit's File Formats (www.vtk.org/pdf/file-formats.pdf) document.

A.1.3 Text

The text format is always in table form, with the dimensions of the table specified in the first line. For example, mesh coordinates can be specified in the following format

```

nno nsd
x1 y1 z1
x2 y2 z2
x3 y3 z3
...

```

Mesh connectivity with

```

nel ndofs
node_1 node_2 node_3 node_4 ...
node_1 node_2 node_3 node_4 ...
node_1 node_2 node_3 node_4 ...
...

```

A generic field with `ndim` components (i.e., scalar, vector, or tensor) is specified by

```
num ndim
f1 f2 f3 ..
f1 f2 f3 ..
...
```

In this last case, the number of rows could refer to either `nno` or `nel`, depending on whether the field is node-based or cell-based.

A.2 Output File Format

The output format for residuals consists simply of a list of scalars for each element in the discretization. If you specify a `.vtk` extension for your output file, this will result in an scalar dataset named `epsilon` stored using the legacy VTK file format in the Cell Data section of the output file. Note that this output consists of the squared values of the local residuals, so further post-processing will be necessary.

Appendix B

License

GNU GENERAL PUBLIC LICENSE Version 2, June 1991. Copyright (C) 1989, 1991 Free Software Foundation, Inc. 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Preamble

The licenses for most software are designed to take away your freedom to share and change it. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change free software – to make sure the software is free for all its users. This General Public License applies to most of the Free Software Foundation's software and to any other program whose authors commit to using it. (Some other Free Software Foundation software is covered by the GNU Library General Public License instead.) You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for this service if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs; and that you know you can do these things.

To protect your rights, we need to make restrictions that forbid anyone to deny you these rights or to ask you to surrender the rights. These restrictions translate to certain responsibilities for you if you distribute copies of the software, or if you modify it.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must give the recipients all the rights that you have. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

We protect your rights with two steps:

1. Copyright the software, and
2. Offer you this license which gives you legal permission to copy, distribute and/or modify the software.

Also, for each author's protection and ours, we want to make certain that everyone understands that there is no warranty for this free software. If the software is modified by someone else and passed on, we want its recipients to know that what they have is not the original, so that any problems introduced by others will not reflect on the original authors' reputations.

Finally, any free program is threatened constantly by software patents. We wish to avoid the danger that redistributors of a free program will individually obtain patent licenses, in effect making the program proprietary. To prevent this, we have made it clear that any patent must be licensed for everyone's free use or not licensed at all.

The precise terms and conditions for copying, distribution and modification follow.

GNU GENERAL PUBLIC LICENSE TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION

0. This License applies to any program or other work which contains a notice placed by the copyright holder saying it may be distributed under the terms of this General Public License. The "Program" below refers to any such program or work, and a "work based on the Program" means either the Program or any derivative work under copyright law: that is to say, a work containing the Program or a portion of it, either verbatim or with modifications and/or translated into another language. (Hereinafter, translation is included without limitation in the term "modification.") Each licensee is addressed as "you."

Activities other than copying, distribution and modification are not covered by this License; they are outside its scope. The act of running the Program is not restricted, and the output from the Program is covered only if its contents constitute a work based on the Program (independent of having been made by running the Program). Whether that is true depends on what the Program does.

1. You may copy and distribute verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice and disclaimer of warranty; keep intact all the notices that refer to this License and to the absence of any warranty; and give any other recipients of the Program a copy of this License along with the Program.

You may charge a fee for the physical act of transferring a copy, and you may at your option offer warranty protection in exchange for a fee.

2. You may modify your copy or copies of the Program or any portion of it, thus forming a work based on the Program, and copy and distribute such modifications or work under the terms of Section 1 above, provided that you also meet all of these conditions:

- (a) You must cause the modified files to carry prominent notices stating that you changed the files and the date of any change.
- (b) You must cause any work that you distribute or publish, that in whole or in part contains or is derived from the Program or any part thereof, to be licensed as a whole at no charge to all third parties under the terms of this License.
- (c) If the modified program normally reads commands interactively when run, you must cause it, when started running for such interactive use in the most ordinary way, to print or display an announcement including an appropriate copyright notice and a notice that there is no warranty (or else, saying that you provide a warranty) and that users may redistribute the program under these conditions, and telling the user how to view a copy of this License. (Exception: if the Program itself is interactive but does not normally print such an announcement, your work based on the Program is not required to print an announcement.)

These requirements apply to the modified work as a whole. If identifiable sections of that work are not derived from the Program, and can be reasonably considered independent and separate works in themselves, then this License, and its terms, do not apply to those sections when you distribute them as separate works. But when you distribute the same sections as part of a whole which is a work based on the Program, the distribution of the whole must be on the terms of this License, whose permissions for other licensees extend to the entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest your rights to work written entirely by you; rather, the intent is to exercise the right to control the distribution of derivative or collective works based on the Program.

In addition, mere aggregation of another work not based on the Program with the Program (or with a work based on the Program) on a volume of a storage or distribution medium does not bring the other work under the scope of this License.

3. You may copy and distribute the Program (or a work based on it, under Section 2) in object code or executable form under the terms of Sections 1 and 2 above provided that you also do one of the following:
 - (a) Accompany it with the complete corresponding machine-readable source code, which must be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
 - (b) Accompany it with a written offer, valid for at least three years, to give any third party, for a charge no more than your cost of physically performing source distribution, a complete machine-readable copy of the corresponding source code, to be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
 - (c) Accompany it with the information you received as to the offer to distribute corresponding source code. (This alternative is allowed only for noncommercial distribution and only if you received the program in object code or executable form with such an offer, in accord with Subsection b above.)

The source code for a work means the preferred form of the work for making modifications to it. For an executable work, complete source code means all the source code for all modules it contains, plus any associated interface definition files, plus the scripts used to control compilation and installation of the executable. However, as a special exception, the source code distributed need not include anything that is normally distributed (in either source or binary form) with the major components (compiler, kernel, and so on) of the operating system on which the executable runs, unless that component itself accompanies the executable.

If distribution of executable or object code is made by offering access to copy from a designated place, then offering equivalent access to copy the source code from the same place counts as distribution of the source code, even though third parties are not compelled to copy the source along with the object code.

4. You may not copy, modify, sublicense, or distribute the Program except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense or distribute the Program is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.
5. You are not required to accept this License, since you have not signed it. However, nothing else grants you permission to modify or distribute the Program or its derivative works. These actions are prohibited by law if you do not accept this License. Therefore, by modifying or distributing the Program (or any work based on the Program), you indicate your acceptance of this License to do so, and all its terms and conditions for copying, distributing or modifying the Program or works based on it.
6. Each time you redistribute the Program (or any work based on the Program), the recipient automatically receives a license from the original licensor to copy, distribute or modify the Program subject to these terms and conditions. You may not impose any further restrictions on the recipients' exercise of the rights granted herein. You are not responsible for enforcing compliance by third parties to this License.
7. If, as a consequence of a court judgment or allegation of patent infringement or for any other reason (not limited to patent issues), conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot distribute so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not distribute the Program at all. For example, if a patent license would not permit royalty-free redistribution of the Program by all those who receive copies directly or indirectly through you, then the only way you could satisfy both it and this License would be to refrain entirely from distribution of the Program.

If any portion of this section is held invalid or unenforceable under any particular circumstance, the balance of the section is intended to apply and the section as a whole is intended to apply in other circumstances.

It is not the purpose of this section to induce you to infringe any patents or other property right claims or to contest validity of any such claims; this section has the sole purpose of protecting the integrity of the free software distribution system, which is implemented by public license practices. Many people have made generous contributions to the wide range of software distributed through that system in reliance on consistent application of that system; it is up to the author/donor to decide if he or she is willing to distribute software through any other system and a licensee cannot impose that choice.

This section is intended to make thoroughly clear what is believed to be a consequence of the rest of this License.

8. If the distribution and/or use of the Program is restricted in certain countries either by patents or by copyrighted interfaces, the original copyright holder who places the Program under this License may add an explicit geographical distribution limitation excluding those countries, so that distribution is permitted only in or among countries not thus excluded. In such case, this License incorporates the limitation as if written in the body of this License.
9. The Free Software Foundation may publish revised and/or new versions of the General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies a version number of this License which applies to it and "any later version," you have the option of following the terms and conditions either of that version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of this License, you may choose any version ever published by the Free Software Foundation.

10. If you wish to incorporate parts of the Program into other free programs whose distribution conditions are different, write to the author to ask for permission. For software which is copyrighted by the Free Software Foundation, write to the Free Software Foundation; we sometimes make exceptions for this. Our decision will be guided by the two goals of preserving the free status of all derivatives of our free software and of promoting the sharing and reuse of software generally.

NO WARRANTY

11. BECAUSE THE PROGRAM IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.
12. IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR REDISTRIBUTE THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively convey the exclusion of warranty; and each file should have at least the "copyright" line and a pointer to where the full notice is found. For example:

One line to give the program's name and a brief idea of what it does. Copyright © (year) (name of author)

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA

Also add information on how to contact you by electronic and paper mail.

If the program is interactive, make it output a short notice like this when it starts in an interactive mode:

Gnomovision version 69, Copyright © year name of author Gnomovision comes with ABSOLUTELY NO WARRANTY; for details type 'show w'. This is free software, and you are welcome to redistribute it under certain conditions; type 'show c' for details.

The hypothetical commands 'show w' and 'show c' should show the appropriate parts of the General Public License. Of course, the commands you use may be called something other than 'show w' and 'show c'; they could even be mouse-clicks or menu items – whatever suits your program.

You should also get your employer (if you work as a programmer) or your school, if any, to sign a "copyright disclaimer" for the program, if necessary. Here is a sample; alter the names:

Yoyodyne, Inc., hereby disclaims all copyright interest in the program 'Gnomovision' (which makes passes at compilers) written by James Hacker.

(signature of Ty Coon)

1 April 1989

Ty Coon, President of Vice

This General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Library General Public License instead of this License.

Bibliography

- [1] Akin, J.E. (2005), *Finite Element Analysis with Error Estimators*, Butterworth-Heinemann, Oxford, 447 pp.
- [2] Encyclopaedia of Cubature Formulas (1998-2005), Katholieke Universiteit Leuven, Dept. Computerwetenschappen, www.cs.kuleuven.be/~nines/research/ecf/
- [3] Hughes, Thomas J.R. (2000), *The Finite Element Method: Linear Static and Dynamic Finite Element Analysis*, Dover, 672 pp.
- [4] Karniadakis, George E. and Spencer J. Sherwin (2005), *Spectral/hp Element Methods for Computational Fluid Dynamics, Second Edition*, Numerical Mathematics and Scientific Computation, Oxford, 657 pp.
- [5] Uesu, D., L. Bavoil, S. Fleishman, J. Shepherd, and C. Silva (2005), Simplification of Unstructured Tetrahedral Meshes by Point-Sampling, *Proceedings of the 2005 International Workshop on Volume Graphics*, 157-165, doi: 10.2312/VG/VG05/157-165.